Lecture 16: Non-blocking caches

Thursday, March 1, 2018 9:40 AM

Outline

- Talk about non-blocking cache controller design
- Talk about Lab 4
- Replacement policies for set associative caches

From last time

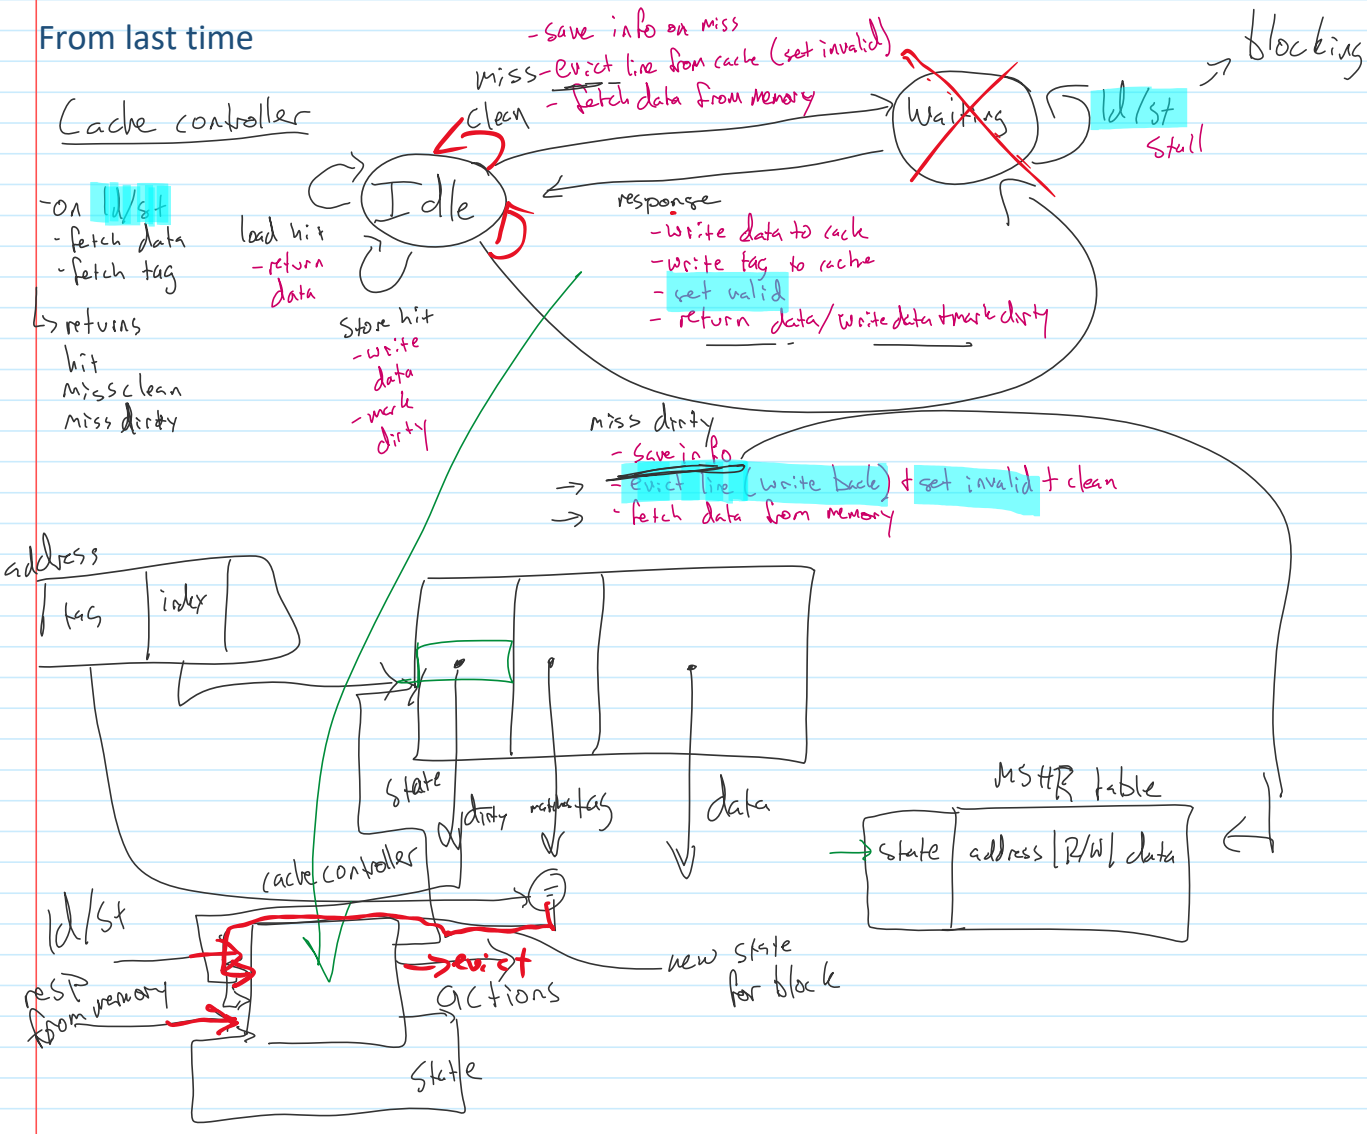

for non blocking cache: Still need to track states
- now tracking state at line/block granularity
- a state machine for every block

→ a way to track outstanding requests

↳ miss status handling register (MSHR)

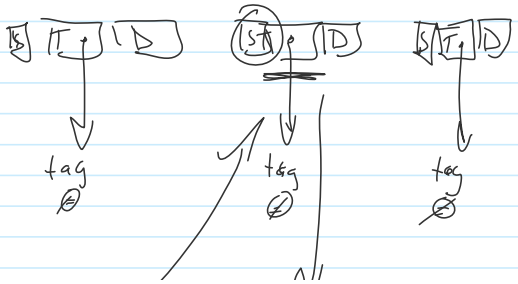

Block-level state machine for a non-blocking

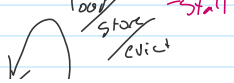

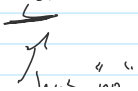

Block-level state machine for a non-blocking write back cache

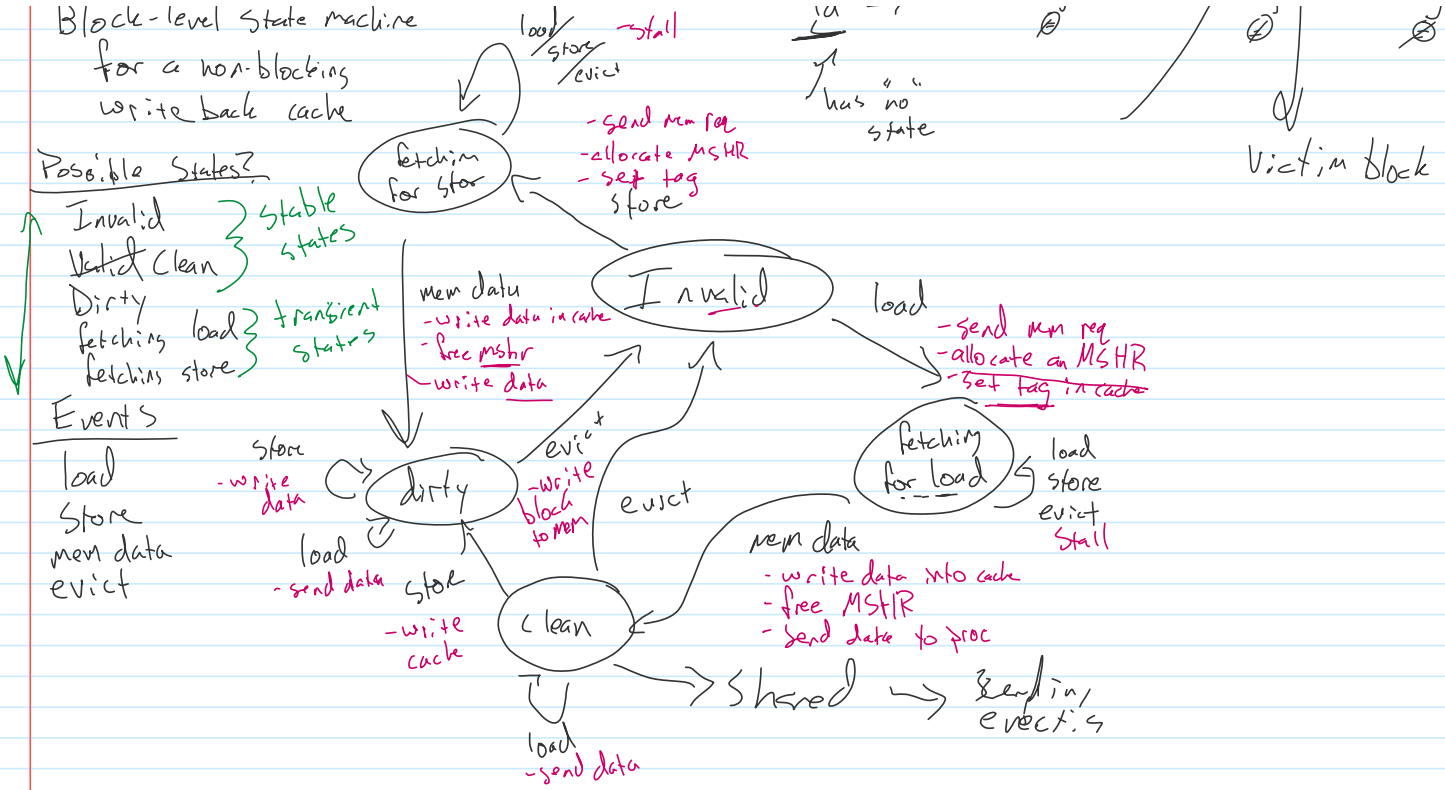

Possible States?

- Invalid
- ~~Valid~~ Clean
- Dirty
- fetching load
- fetching store

(Invalid, Clean: stable states; fetching load, fetching store: transient states)

Events

- load
- store
- mem data
- evict

Replacement policies

How to choose what to evict?

Least recently used (LRU) evict the item used furthest in past

+ temporal locality
+ Close to theoretical best

- expensive in bits, and hardware
- every access is expensive

round robin way 0, way 1 ...

Clock algorithm → give things a "time"

Random

not most recently used NMRU ←

Vary w/ level of cache