

Lecture 19: Parallel architectures

Tuesday, March 13, 2018 10:46 AM

Outline

- Types of parallel architectures
- Taxonomy for parallel architectures
- Emerging parallel architectures

Why parallelism?

Incr throughput w/ multiple tasks

Incr performance/reduce latency

2 choices \rightarrow double freq \rightarrow w/w limits
or double cores \rightarrow lower power

$$\rightarrow P = \frac{1}{2} C V^2 f$$

- Scalar \rightarrow not parallel single inst stream + single data stream

Types of parallel architectures

\rightarrow

how to program?

how to communicate?

how are instructions executed?

(GPU)
★ - Accelerators + CPU

\rightarrow send a task to accelerator

\rightarrow send data to accel explicitly

\rightarrow offload computation

\rightarrow explicitly copy results back

★ Multi core

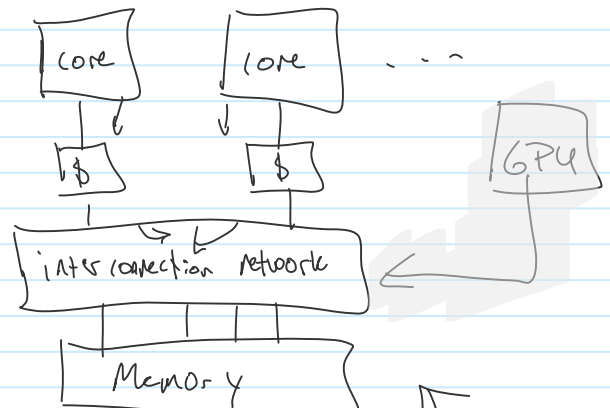
\rightarrow Shared memory

\hookrightarrow because all processors see same data at same address

\rightarrow reading/writing memory for comm.

\rightarrow programming?

threads \rightarrow has a program counter
 \hookrightarrow ...



Announcements

- course evaluations

- lab 5 out

\hookrightarrow yes Partner

\hookrightarrow **no** late submissions

- Final Wellman '26

Review Friday during discussion

\rightarrow office hours Mon 4:30-6

→ programming?

threads → has a program counter
but in the same virtual
address space
usually executing same code

Single program multiple data (SPMD)

SPMD is a type of MIMD system

↳ multiple inst. multiple data

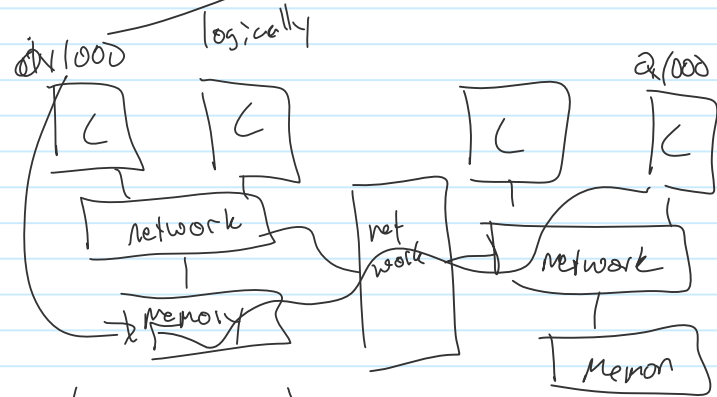
★ Multi-processor / multi-chip

has multiple processor "sockets"

Shared memory

- memory access time is non-uniform

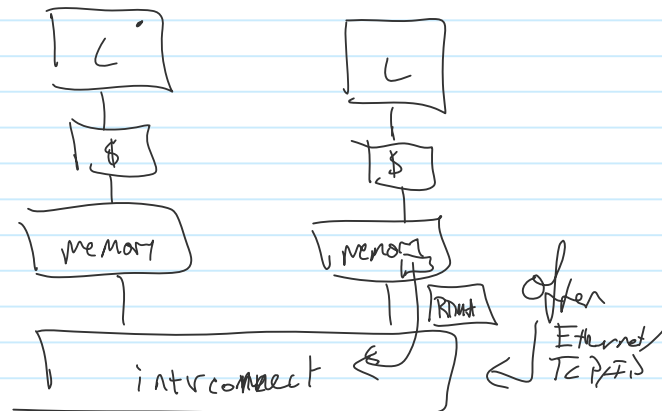
- programmers must think about data
locality



★ Warehouse-scale computers (cloud / datacenter / super computer)

- communicate with message passing
request for data
respond w/ data

- remote DMA (Direct memory access)



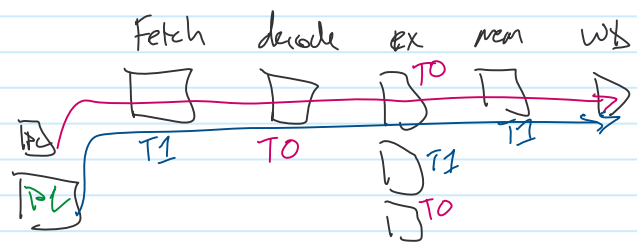
★ Multi-threaded process (Intel hyperthreading)

4 cores / 8 threads

→ increases throughput
at cost of latency

SPMD

shared memory



★ - vector architectures

→ take advantage of data parallelism

$$\vec{a} + \vec{b} = \vec{c}$$

... for loop

→ take advantage of **data parallelism**

$$\vec{a} + \vec{b} = \vec{c}$$

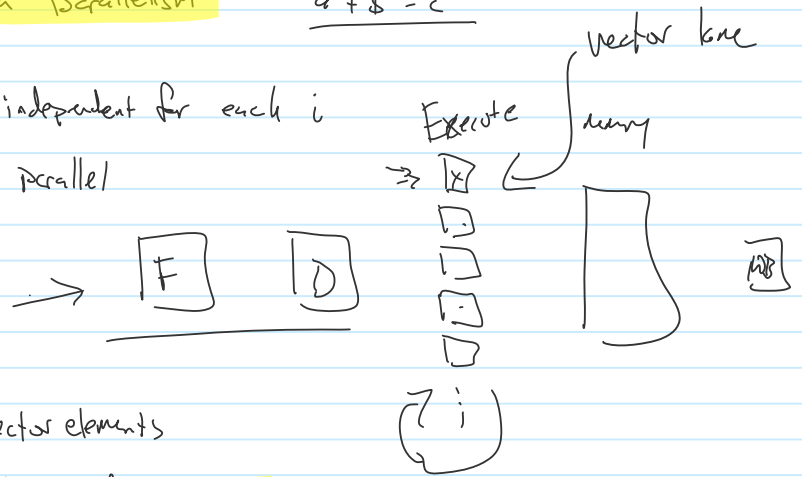
for ($i = 0 \rightarrow n$)
 → $c[i] = a[i] + b[i]$ independent for each i

execute all vector elements in parallel

- incr throughput & perf.
- easy to reason about

- fetch/decode shared
 decode once execute for all vector elements

- **Single instruction stream multiple data streams**
 SIMD



		Inst. Stream	
data stream	Single	Single	multiple
	Single	SIMD scalar	MISD? not really... MIMD SPMD + threading shared memory message passing
multiple		SIMD vectors/ GPUs	

Examples of parallel programs

```
void daxpy(double *X, double *Y, double alpha, const int N)
{
    for (int i = 0; i < N; i++)
    {
        Y[i] = alpha * X[i] + Y[i];
    }
}
```

scalar

```
void daxpy_simd(double *X, double *Y, double alpha, const int N)
{
    __m256d *vx = (__m256d*)X;
    __m256d *vy = (__m256d*)Y;
    __m256d va = {alpha, alpha, alpha, alpha};

    for (int i = 0; i < N/4; i++)
    {
        __m256d mult = _mm256_mul_pd(vx[i], va);
        __m256d add = _mm256_add_pd(vy[i], mult);
        vy[i] = add;
    }
}
```

```

}

void daxpy_fmadd(double *X, double *Y, double alpha, const int N)
{
    __m256d *vx = (__m256d*)X;
    __m256d *vy = (__m256d*)Y;
    __m256d va = {alpha, alpha, alpha, alpha};

    for (int i = 0; i < N/4; i++)
    {
        vy[i] = _mm256_fmadd_pd(va, vx[i], vy[i]);
    }
}

void daxpy_threads(double *X, double *Y, double alpha, const int N)
{
    int num_threads = std::thread::hardware_concurrency();

    std::thread threads[num_threads];

    int chunk = N / num_threads;

    for (int i=0; i<num_threads; i++) {
        threads[i] = std::thread(daxpy, X+i*chunk, Y+i*chunk, alpha, chunk); ←
    }

    for (int i=0; i<num_threads; i++) {
        threads[i].join();
    }
}

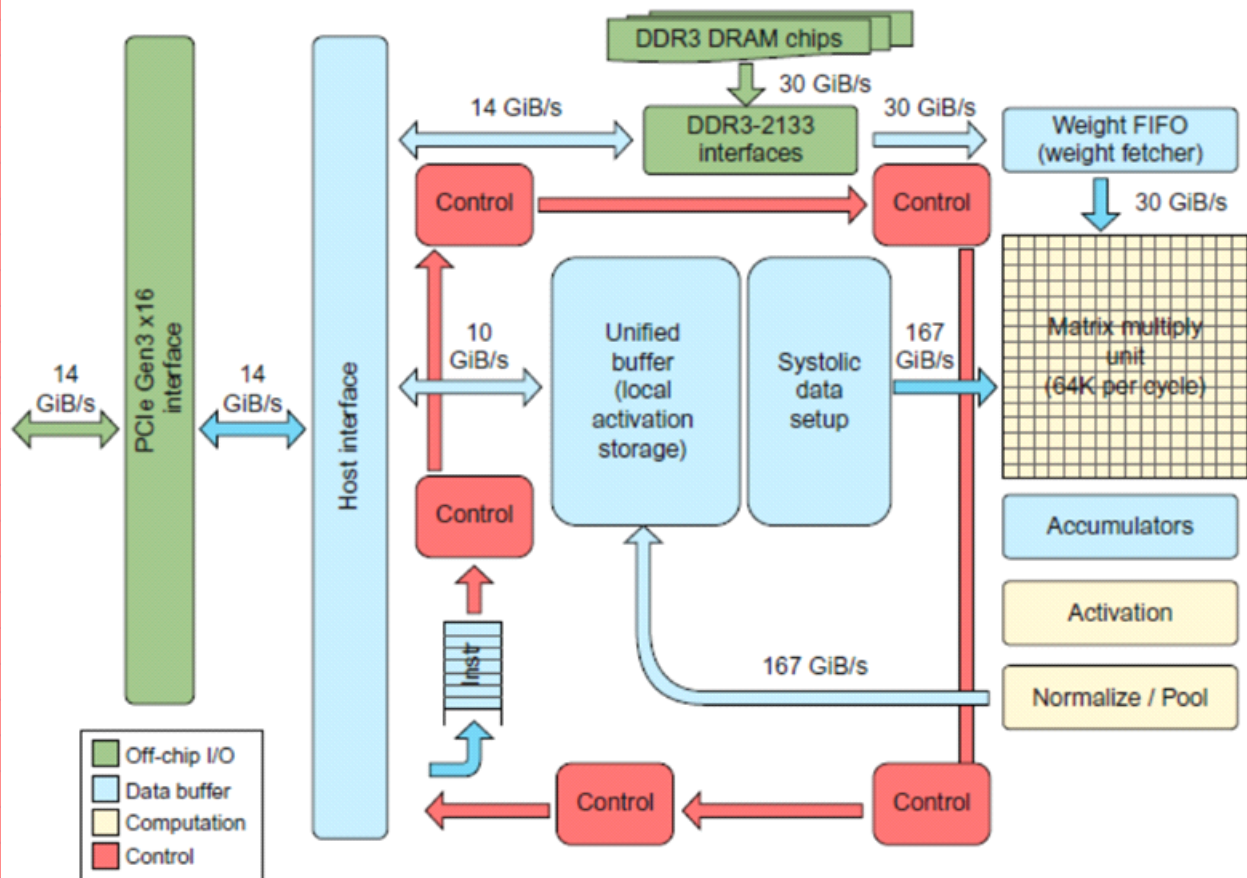
void daxpy_omp(double *X, double *Y, double alpha, const int N)
{
    #pragma omp parallel for
    for (int i = 0; i < N; i++)
    {
        Y[i] = alpha * X[i] + Y[i];
    }
}

void daxpy_opencl(double *X, double *Y, double alpha, const int N)
{
    int tid = get_global_id(0);
    for (int i = tid; i < N; i += get_global_size(0)) {
        Y[i] = alpha * X[i] + Y[i];
    }
}

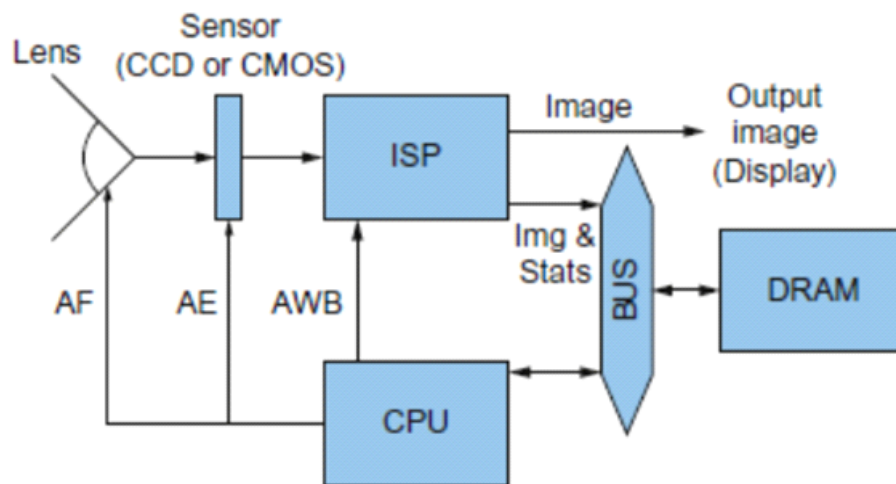
```

New parallel architectures

Google's TPU



Pixel visual core



5 x 5 stencil

